

Software engineering tools for OP2 project

Mike Giles

September 14, 2010

Abstract

As we begin phase 2 of the OP2 project, with an increased amount of collaboration between different developers, I think it is a good time to consider the software engineering tools we will use.

As an Associate Director of the Oxford e-Research Centre, I'm also interested in this topic to identify and promote best practice in software projects within OeRC.

This is an area in which I am aware that my own knowledge and practice is deficient, so many of my suggestions are based on the recommendations of others, and I will be interested to hear people's opinions on them.

Tools

- Version control

I've been told that [git](#) and the web-based [GitHub](#) are much better than the older [CVS](#) and [SVN](#).

- Make system

Currently, we use GNU makefiles, but in the future we should consider using [CMake](#) which is platform independent and has [CUDA support](#).

- Compilers

Currently we are using NVIDIA's [nvcc](#) for the CUDA code, and [g++](#) for the C/C++ code, but in the future we will be using PGI's [CUDA FORTRAN](#) compiler, and also ought to consider [Intel's compilers](#) for multithreaded C/C++/FORTRAN code and to produce AVX executables.

- MPI

There are several different MPI implementations. We will probably continue to use the one recommended / supplied by the hardware supplier.

For parallel file I/O, we plan to use [Parallel HDF5](#) which is [built on MPI-IO](#), rather than directly using [MPI-IO](#).

- Parallel debugging

I think Allinea's [DDT](#) is the best choice for both MPI and CUDA applications. At present the [CUDA support](#) is only for CUDA version 3.0, but 3.1/3.2 should be supported in the near future

- Performance monitoring and optimisation

I think Allinea's [OPT](#) may be good for MPI applications, but I don't know if it has any support for CUDA.

Stephen Jarvis' group models the performance of parallel applications using network data obtained from [Scalasca](#) and low-level CPU performance data from [PAPI](#).

Another possibility for low-level optimisation of CPU code (particularly multithreaded code?) is Intel's [vtune](#) software.

- Code documentation

At present, I am happy to write documentation in LaTeX, and that allows me to write different sets of documentation for users and developers.

However, some people recommend the use of [Doxygen](#) and this could be considered to generate the definitive developer reference.

- Build-and-test

I'm told that [Buildbot](#) is the standard choice for automated compilation and testing.

- Coverage testing

This is something I've never used before, but tools such as [gcov](#) and [lcov](#) check which bits of your code are being exercised by the testcases used to validate the code.

- Code generation

Currently, the OP2 code generator is written in MATLAB. This is not an ideal choice; it's simply a language I am comfortable with. Paul Kelly's group may re-write it entirely in [rose](#) or it could be ported to python.