OPlus2 for many-core platforms

Mike Giles

February 10, 2010

Abstract

This document lays out the plans for a new open source OPlus2 library for unstructured grid parallel execution on NVIDIA GPUs and other many-core hardware platforms.

This document will be all-encompassing to begin with. Later, it will make sense to split it into different documents (e.g. Users Guide; Developers Guide; Algorithms for partitioning, colouring and halo construction) but for now I think it is simplest to have it all together in one place.

1 Motivation

The original OPlus library was developed more than 10 years ago for distributed memory parallel execution of algorithms based on unstructured grids. It was originally developed in PVM by Paul Crumpton and myself, but more recently Nick Hills ported it to MPI and made numerous internal changes (such as parallel partitioning) which greatly improved its performance.

The strengths and weaknesses of the existing OPlus are discussed below, but the main motivation for a new version is the emergence of multicore architectures. Not only are Intel, AMD, and others all pursuing a multicore CPU strategy, with up to 6 cores today but maybe 16 cores (and up to 64 threads?) within a few years, but also NVIDIA and ATI (now part of AMD) have graphics cards with huge numbers of cores. The current NVIDIA GTX280 GPU, for example, has 240 cores, each of which is as capable as an Intel Core 2 Duo core in single precision.

The objective of OPlus2 is to provide a library which will enable efficient execution on these new architectures at the lowest level, while retaining the ability to perform distributed memory parallelisation on the highest level.

2 OPlus

After more than 10 years of use, it is a good time to review the strengths and weaknesses of the existing OPlus library.

Strengths:

- general distributed set/pointer representation has worked well

- single application source code leading to both sequential and parallel executables

- fairly clean definition of parallel loops

- good scalability on big problems with up to 10^8 grid nodes on 100+ processors (partly due to the original design but largely due to Nick Hills' improvements)

- works well under both Unix/Linux and Windows operating systems

Weaknesses:

- file I/O is poor, not parallel

- `OP_ACCESS` declaration of data dependencies in parallel loops is tedious, and errors can be hard to debug

- only FORTRAN is supported, and it is only used by the HYDRA CFD code

The natural goal with OPlus2 is to retain the strengths, address the weaknesses, and add the capability to efficiently exploit new many-core architectures.

References

1) P.I. Crumpton and M.B. Giles. 'Multigrid aircraft computations using the OPlus parallel library'. in Parallel Computational Fluid Dynamics: Implementations and Results Using Parallel Computers, 339-346. A. Ecer, J. Periaux, N. Satofuka, and S. Taylor, editors, North-Holland, 1996. (PDF)

2) D.A. Burgess, P.I. Crumpton, and M.B. Giles. 'A parallel framework for unstructured grid solvers'. In Computational Fluid Dynamics '94: Proceedings of the Second European Computational Fluid Dynamics Conference, pages 391-396. S. Wagner, E.H. Hirschel, J. Periaux, and R. Piva, editors. John Wiley and Sons, 1994. (PDF)

3 Key assumptions and requirements

The key assumptions are the same as for OPlus: that the application involves a number of static sets (e.g. nodes, edges, faces, cells) connected by pointers, and each step in the algorithm involves an operation applied to all elements of a set (with at most one level of pointer indirection to data associated with another set) and the result is independent of the order in which the elements are processed (e.g. Jacobi iteration is allowed but not the standard Gauss-Seidel).

The requirements are:

- single application source code leading to efficient execution on

    - multicore CPUs using OpenMP and/or SSE/AVX instructions
    - GPUs using CUDA or OpenCL

- very good scalability up to very large problem sizes using MPI to exploit multiple GPUs and/or multicore CPUs within a cluster

- support for both C/C++ and FORTRAN

- support for Unix/Linux and Windows operating systems

- options for extensive error-checking

- compatibility with HYDRA's use of automatic differentiation for generating linear and adjoint code

With OPlus, the user code was compiled and then linked to either a sequential single-thread library or a parallel library (based on MPI message passing). There was no use of program transformation, partly because it wasn't needed and partly because we had no knowledge of such tools.

With OPlus2, the efficient support of multicore CPUs as well as GPUs will require the use of program transformation tools, with the expertise being supplied by Paul Kelly's group at Imperial College. However, I think it is still desirable to have a non-efficient single thread CPU implementation which does not require program transformation; this can be used for application development and debugging and should include the extensive error-checking tools mentioned above.

4 Outline workplan

The development of OPlus2 will proceed in two phases, with perhaps some overlap between the two.

- Phase 1 will develop OPlus2 for use on a single GPU or a single shared-memory system with one or more multicore CPUs (using OpenMP and SSE/AVX instructions).

- Phase 2 will extend this by adding on top an MPI layer which is very similar to the existing OPlus implementation. This will enable distributed-memory computations on a cluster of nodes, each with one or more GPUs and/or multi-core CPUs.

The partitioning and parallelisation strategies required by the two phases will be similar in some ways, but will also be different because of differences in key hardware aspects. For example, typical cluster compute nodes have GBs of shared memory, whereas each "multiprocessor" in an NVIDIA GPU currently has only 16kB of shared memory. As explained later, this leads to a different approach in handling data dependency constraints.

The phase 1 workplan is as follows (IC = Imperial College post-doc):

- single-thread CPU version of test code written in C, including non-efficient CPU library (Mike, Sept 09)

- CUDA version of test code written in C (Mike, Oct/Nov 09)

- document describing program transformation requirements for C → CUDA (Mike/Paul, Dec 09)

- single-thread CPU version of test code written in FORTRAN, including non-efficient CPU library (Mike, Jan 10)

- C and FORTRAN versions of 2D airfoil test code (Mike, Mar 10)

- program transformation for C → CUDA (IC, June 10)

- OpenMP/SSE version of original test code (Mike, June 10)

- program transformation for FORTRAN → CUDA (IC, Sept 10)

- test program transformation tools on 2D airfoil test code (Mike, Sept 10)

- program transformation for C → OpenMP/SSE parallel loops (IC, Dec 10)

- OpenCL version of original test code (Mike, Dec 10)

- "lint"-style syntactic analysis and error-checking for parallel loop kernels (IC, Mar 11)

- program transformation tool to aid porting of HYDRA OPlus calls to OPlus2 (IC, June 11)

- begin phase 2 (Mike, July–Sept 11)

- program transformation for C, FORTRAN → OpenCL (IC, Sept 11)

At the end of the first phase we will have a viable library and program transformation tools for execution on a single host with one or more NVIDIA GPUs.

The porting of HYDRA will overlap with the final parts of phase 1 and the early parts of phase 2, with the three main steps being:

- convert remaining parallel loops into kernel format (this can be done at any time, independent of work on phase 1);

- convert OPlus parallel loops into OPlus2 parallel loops, aided by the tool developed by Imperial;

- convert file I/O to the new parallel file I/O agreed for phase 2.

The timing for phase 2 is a bit uncertain, but the key steps (some of which can be overlapped) are:

- review existing OPlus algorithms and code (Mike, 3 months?)

- review HDF5 and MPI I/O and design OPlus2 file I/O (Leigh, 3 months?)

- review key implementation issues and concerns (Nick, 3 months?)

- implement new file I/O (1 month?)

- implement phase 2 data partitioning and halo construction (3 months?)

- test on a conventional cluster (1 month?)

- optimise GPU partitioning to overlap MPI communication with GPU computation (1 month?)

- test on a GPU cluster (1 month?)

- test HYDRA on a GPU cluster (1 month?)

- develop performance monitoring tools to guide further development (1 month?)

The work in most of these items will be done jointly by the Imperial College postdoc and myself.

5 Data access and dependency conflicts

Suppose that one wants to solve the sparse system of linear equations

$$A\,u + r = 0$$

using Jacobi iteration with

$$u^{n+1} = u^n + (A\,u^n + r)\,.$$

A can be stored in a sparse edge-based format `A,p1,p2` in which each non-zero element has a corresponding "edge" `e` with `A[e]`$\equiv A_{i,j}$, `p1[e]`$\equiv i$, `p2[e]`$\equiv j$.

Using this, the key steps in the algorithm can be implemented using the C routines in Figure 1. The access patterns for the arrays can be classified as

- "r" read-only

- "w" write-only

- "b" both read and write

- "i" increment (a special case of read/write)

In `res`, the access for `A, u, du` is "r", "r" and "i", respectively, while in `update`, the access for `u, du, r` is "i", "b" and "r", respectively.

It is clear from looking at these routines that it makes no difference the order in which the edges are processed in `res`, or the nodes are processed in `update`, and so these conform to the requirements for OPlus parallelisation. The one potential

```
void res(int ne, int *p1, int *p2, float *A, float *u, float *du) {
   for (int e=0; e<nedge; e++)
      du[p1[e]] += A[e] * u[p2[e]];
}

void update(int nn, float *u, float *du, float *r) {
   for (int n=0; n<nn; n++) {
      u[n]  += du[n] + r[n];
      du[n] = 0.0f;
   }
}
```

Figure 1: C routines for simple Jacobi iteration

problem is a potential data dependency conflict in incrementing the array `du` in `res`. Because of the indirect addressing using `p1`, it is possible that different edges may try to update the same element of `du` "at the same time". Unless the hardware supports atomic updates for floating point variables, this could lead to incorrect results.

In OPlus, the parallelisation partitions the data so that each partition "owns" some of the edges and some of the nodes. Each partition "executes" the edges and nodes which affect the data it "owns". Near partition boundaries, this may require one partition to execute an edge belonging to a different partition. In addition, if the execution of the edge requires the use of data belonging to another partition, a copy of that data has to be made; this is referred to as "halo" data in the distributed-computing community.

Within each partition, a single thread sequential execution model is used and so there is problem with the data dependency issue. However, in more complex applications where executing an edge may involve updates to more than one node, there is a redundancy/duplication issue, since the same edge may be executed by more than one partition. This is acceptable provided the partition sizes are large enough so that the number of halo edges and nodes is relatively small.

In OPlus2, the same strategy will be used in phase 2 for distributed computing across multiple GPUs, but in phase 1 the same approach will not be acceptable because the partition sizes to fit into each of the multiprocessors within a single GPU are so small that there will be a significant level of duplication. We will instead use an alternative partition "colouring" strategy which is described in a later section.

6   CUDA parallelisation strategy

The first step is to reorder/renumber all sets to improve data locality, so that neighbouring nodes or edges have indices, and therefore memory locations, which are close to each other. There are lots of ways in which this might be done, but one simple effective way which we have used in the past is recursive coordinate bisection. For a 3D dataset, one would first bisect each set in the x-direction, then the y-direction, then the z-direction, and then repeat the sequence recursively to build up a binary tree whose leaves are then numbered sequentially.

The idea then is to parallelise loop operations by using blocks (mini-partitions) which are small enough to fit into the limited shared memory on each multiprocessor within an NVIDIA GPU. For the edge operations in `res` in Figure 1, the operations for each block would be

- load in required values of `u` into shared memory

- calculate increment

- add increment to shared memory array `inc`

- add increment `inc` to array `du` in main graphics memory

The first step of loading `u` into shared memory is performed because the same values of `u` are likely to be used by several edges, and so doing this minimises the total data transfer requirements from the main graphics memory.

The second step is simple, and can be performed in parallel by all threads in the block (with just one edge per thread?) without any data conflicts. The potential data conflict is in the third step when the increments are combined. This will be avoided by "colouring" the threads so that no two threads update the same variable at the same time. This is a well-established technique for avoiding data conflicts in vector code.

(To minimise the number of colours within each thread "warp", it may be best to reorder the threads within each block. This would scramble up the identity mapping, used here for the array `A`, so there would be some cost but I think there would be savings overall.)

The fourth step adds the combined increments to the data in the main graphics memory. Here again there is a potential data conflict problem due to different blocks trying to update the same data. This will be avoided by colouring the blocks so that no two blocks of the same colour update the same data. The blocks would be executed one colour at a time, with a synchronisation between each colour.

6.1   Execution plans

In standard MPI computations, the partitioning is done once and the same partitioning is used for stage of the computation. This is because the cost of re-partitioning the data is excessive.

In contrast, between each stage of the computation on the GPU the data resides in the main graphics memory, and so the blocking for each loop calculation can be considered independently of the requirements of the other parallel loops.

Based on ideas from FFTW, I plan to construct for each parallel loop a "plan", which is a customised blocking of the execution on the GPU for that loop, making optimum use of the local shared memory on each multiprocessor considering in detail the memory requirements of the loop computation.

The plan will also include the colouring of the blocks, and colouring of the threads within a block, to avoid data conflicts. A typical CFD calculation may need 50-200 plans each of which may be executed more than once during a single timestep or multigrid iteration.

6.2   Some rough numbers

A big HYDRA CFD calculation will use almost all of the 4GB on a Tesla card. Given that the local shared-memory size is 16kB, this suggests that the data intensive parallel loops will have over 100,000 blocks. [1]

10 colours might be needed to avoid data conflicts, suggesting up to 10,000 blocks per colour. Given this number it doesn't matter if there are variations in the amount of compute per block, the important thing is to make sure each block uses as much shared memory as possible.

HYDRA needs up to 40 floats per grid node. 16kB corresponds to 4000 floats which is equivalent to 100 nodes, which is roughly 5^3. This should be big enough to get a fair degree of re-use of nodal data within the block, maybe 50% of maximum. The block may have 400 edges, with maybe 40 per colour so we need to use colour-locked increments, with thread synchronisation in between colours.

---

[1] On the one hand, this will be an underestimate since some of the data will be used by more than one block, but on the other hand it will be an overestimate because not all of the data will be needed for just one loop.

7 OP2 execution

An OPlus2 calculation has three stages of execution, each with a number of steps. In the description below, the steps in parentheses are the extra ones needed for the phase 2 MPI parallelisation:

1. initialisation

   - declare sets, pointers and datasets
   - (partition sets and renumber pointers)
   - (compute halos)
   - (load and distribute datasets)
   - renumber elements and pointers
   - initialise the GPU, copy constants to its constant memory and copy datasets to its global memory

2. parallel computation

   - (exchange halo data as necessary)
   - build new CUDA execution "plan", if necessary
   - execute "plan"

3. termination

   - copy datasets back to host
   - terminate CUDA operations on the GPU
   - (write datasets to disk)

8 Parallel loop syntax

My initial idea is that the parallel loop should have the following syntax:

op_par_loop(kernel, set, label,
$\qquad$ **arg1, ptr1, index1, type1, dim1, access1,**
$\qquad$ **arg2, ptr2, index2, type2, dim2, access2,**
$\qquad$ **..., ..., ..., ..., ..., ...,**
$\qquad$ **argN, ptrN, indexN, typeN, dimN, accessN)**

kernel	name of the kernel function to be applied to each element
label	a label string for diagnostic purposes
set	set whose elements are to be processed
arg1	first dataset
ptr1	pointer from main set to set on which **arg1** is defined
index1	index of pointer to access **arg1**
type1	type of dataset (e.g. "float", "int")
dim11	dimension of dataset
access1	information on how **arg1** is used within kernel "r" = read-only "w" = write-only "b" = both read + write "i" = increment

Here **kernel** is a routine which process a single element in **set**. This will get converted by a pre-processor into a routine called by the CUDA kernel function. The pre-processor will also take the specification of the arguments and turn this into the CUDA kernel function which loads in the data from the device main memory into the shared storage, then calls the converted **kernel** for each element, then writes back the updated data to the device main memory.

Parallel loops of the above form will form the main bulk of the application code, after an initialisation stage in which the sets, the pointers between the sets and the data associated with the sets are all declared.

This syntax combines the functions of **OP_PAR_LOOP** and **OP_ACCESS** in the current OPlus, and the specification of the single element **kernel** function fits in well with the use of automatic differentiation in HYDRA.

There is a fair amount of redundant information here, but the implementation will check it is consistent; see discussion in Blog section.

There are also routines to declare sets, pointers and datasets:

op_decl_set(elems, dim, coords, set, label, mapping, maplabel)

elems	number of elements in the set
dim	dimension of the coordinate data (for renumbering)
coords	coordinate data
set	structure containing set info (set by routine)
label	a label string for diagnostic purposes
mapping	an identity mapping (needed for some parallel loops)
maplabel	a label for the identity mapping

op_decl_ptr(from_set, to_set, dim, ptrdata, ptr, label)

from_set	set from which it points
to_set	set to which it points
dim	number of pointers per element
ptrdata	input pointer data
ptr	structure containing ptr info (set by routine)
label	a label string for diagnostic purposes

op_decl_dat(set, dim, type, data, dat, label)

set	set with which the data is associated
dim	amount of data per element
type	type of data (e.g. "float", "int")
data	input data array
ptr	structure containing data info (set by routine)
label	a label string for diagnostic purposes

9   Execution plan construction

New notes: 1/11/09

I've almost finished the coding for the execution plan construction so it's time to make some notes, mainly for my own benefit but also to help anyone reading the code.

A call to **op_par_loop** leads to a call to **plan** which looks to see if there is an existing plan. If there is, it returns its ID, otherwise it constructs a new plan and returns its ID. There is a match to an existing plan if all of the arguments match. Note that this needs to be a run-time match; it is not something that can be decided at compile-time since the same call in the code can be used for different levels within a multigrid solver, and each level requires its own execution plan.

In constructing a new plan, the first step is to see if the loop involves any indirection. If it does not, the execution plan and the corresponding CUDA code is very simple.

When there is indirection, the construction of the plan involves the following steps:

- identification of the datasets involving indirection, and the "increment" sub-set of these which involve an **OP_INC** access producing a potential data conflict which requires coloring

- for each thread block:

  - for each indirection dataset:
    * assemble list of elements pointed to
    * sort into ascending order and eleminate duplicates
    * store away for use by CUDA routine
    * re-number pointers accordingly and store them

  - color the threads (and perhaps re-order threads to group by color to minimise warp divergence?)

- color the blocks (and perhaps re-order the blocks to group by color?)

9.1 Coloring

Coloring is potentially time-consuming, so an important "trick" here is to assign 32 colors at a time using bit operations. I explain it here for thread coloring but the same approach is also used for block coloring.

The key to the implementation is an array with a single 32-bit unsigned integer for each indirectly addressed increment set element. This is initialised to zero and then, as colors are assigned, the corresponding bits will be set to indicate the color of the elements accessing that indirectly addressed element.

The algorithm proceeds as follows:

- loop over all elements which have not yet been colored, and for each one

  - initialise mask variable to zero
  - loop over all indirect increment set elements and "add" to the mask using the logical OR operation
  - find first bit of mask which is <u>not</u> set; this is the color of the new element (provided there is a bit not set)
  - reset the mask to have the appropriate color bit set
  - loop back over all indirect increment set elements "adding" the mask usign the logical OR operation to set the appropriate the color bit

- if 32 colors weren't sufficient, reset the work array to zero and repeat the whole procedure for those which have not yet been colored

In most cases I expect that 32 colors will be sufficient, and so a single pass through each thread block will be all that is required for the thread coloring, and then a single pass through the entire set for the block coloring.

9.2 Different access type for same dataset

What should be done if different argument involve the same dataset with different access types?

- treat each access type as giving a different indirect dataset for the purposes of the execution plan construction

- requires there to be <u>no</u> data conflict between these different sets; must be explained carefully in documentation, and checked in code which validates correct usage

10 Blog

Here I make various quick notes, partly for my own benefit so that I can remember why I did something, or to record things for the future, and partly to inform others about the reasons for my actions.

The thoughts are not carefully organised or well explained – please ask questions if you think something looks interesting but unclear.

- In designing the parallel loop syntax I wanted something simple and clean, without lots of redundant information. This is how I designed the initial CPU implementation, but the more I think about it and how the source code transformation will turn it into optimised code, the more I think it's better to go with something closer to the original OPlus syntax with redundant information. This will enable compile-time optimisations (e.g. unrolling of loops copying data into local shared-memory arrays), and the redundant info can be checked for consistency.

- I'm not sure how best to handle different datatypes. My initial implementation supported only datasets composed of floats. Rather than having different dataset types for each basic type, I think it may be best to change the struct to use void pointers and store the underlying datatype. Alternatively, can a C++ guru construct a parameterised class? If I use void pointers, then I should do run-time type-checking.

- I'm going to have to be careful about the total size of all of the execution plans, but I think they won't be too large compared to the datasets.

- The optimal block sizes may be hard to determine a priori. It may make sense to use dynamic optimisation, adjusting them at run-time to reduce execution time.

- When pointers are not used, data can be loaded straight into registers, not shared-memory.

- When pointers are used, it's best to use shared-memory rather than relying on caching because with shared-memory only need to store the data which is needed.

- In the future, the labels for various routines can be handled automatically either through using macros or through the source code transformation.

- There are alternatives to the thread-colouring approach. One is to use atomic locks, and another is to store the increments in shared-memory without combination then use one thread for each output to combine things sequentially. Ideally, it would be good to try each of these alternatives.

- Coping with a variable number of parameters might be painful. C++ has `va_list` to handle a variable number of inputs, but unlike MATLAB I think there's no support for calling a function with a variable number of parameters. FORTRAN can't do either which is why OPlus had a separate call for each argument, which coincidentally is similar to what happens in OpenCL. Maybe I should stick with that approach for FORTRAN, and for C++ use a template with the number of parameters being specified as a template parameter?

- The `m` loops in `op_par_loop` would be unrolled by the source code transformation, and lots of code eliminated by evaluating many of the `if` tests. The NVIDIA compiler would then automatically unroll the little `p` loops in most cases.

11   Preprocessor

I have developed a prototype preprocessor in MATLAB. This parses the input source code, looking for all `op_par_loop` calls. It then parses their arguments and based on this it generates a CUDA kernel routine which will execute on the GPU, and a host stub code which will select the execution plan and call the kernel routine. Each of these aspects will now be described in some detail.

11.1   Argument parsing

The arguments are parsed to identify which datasets are addressed indirectly, using pointers, and which are addressed directly.

Those which are addressed indirectly are further analysed to see if the same indirect dataset is used for more than one argument with the same access type (e.g. read-only, or increment) but perhaps a different pointer (or a different index within a vector pointer table). For example, in the HYDRA CFD code, a flux calculation for a single edge has 2 arguments which correspond to the flow variables at either end of the edge. It is the same flow variable dataset in each case, but the first edge-node pointer index is used for the first node, and the second is used for the second node.

Identifying shared indirect datasets is important to minimise the use of shared memory within the GPU.

At the end of the parsing process, the preprocessor has a list of all indirect datasets and their access types, and a list of all of the main arguments and their access types and also their mapping to indirect datasets where relevant.

There is an important assumption that if the same dataset is used for different arguments with different access types, then there is no data conflict between the two. Therefore, each will be treated by the parser and the subsequent code generated by the preprocessor as a distinct indirect dataset.

11.2   Kernel code generation

If there are no indirect datasets, the kernel code which is generated is relatively simple. Each thread block executes one section of the set. Read-only data is read from global memory into local registers, and increment arguments are initialised to zero. A suitably-wrapped copy of the user's kernel function is then invoked as an inlined routine. Write-only data is written back to the global memory, and for increment data the values held in global memory are incremented.

If there are indirect datasets then the kernel code is significantly more complicated. Each thread block again executes one section of the set, but now shared

memory is used to hold the indirect datasets. The first step is to read in the appropriate read-only indirect datasets and initialise to zero the increment indirect datasets. Next comes the execution phase, in which the indirect dataset values are taken from the shared memory arrays. Thread coloring is used during the incrementing of the indirect dataset values, to prevent data conflict. Finally, after the completion of the execution phase the write-only indirect datasets are written to global memory, and the global memory values for the increment indirect datasets are incremented accordingly.

For further details it is best to read the code which is generated complete with comments describing each phase of its execution.

11.3   Stub code generation

If there are no indirect datasets, the stub code simply calls the kernel routine once with the required number of thread blocks.

If there are indirect datasets, the stub code defines the mapping from arguments to indirect datasets and then calls a library routine which either calculates an appropriate execution plan, or identifies one which has been generated previously using the same inout arguments. The stub code then calls the kernel routine several times, once for each block color required by the execution plan to avoid data conflicts between the blocks.

11.4   New main code generation

The final task for the preprocessor is to create a new main code. This involves changing the header file which is used, defining a number of routine prototypes, and changing slightly the syntax of the `op_par_loop` calls.

11.5   Build process

Figure 2 shows the build process for a single thread CPU executable. The user's main program (in this case `jac.cpp`) uses the OP header file `op_seq.h` and is linked to the OP routines in `op_seq.c` using `g++`, perhaps controlled by a Makefile.

Figure 3 shows the build process for the corresponding CUDA executable. The preprocessor parses the user's main program amd produces a modified main program and a CUDA file for each of the kernel functions. These are then all compiled and linked to the OP routines in `op_lib.cu` using `g++` and the NVIDIA CUDA compiler `nvcc`, again perhaps controlled by a Makefile.

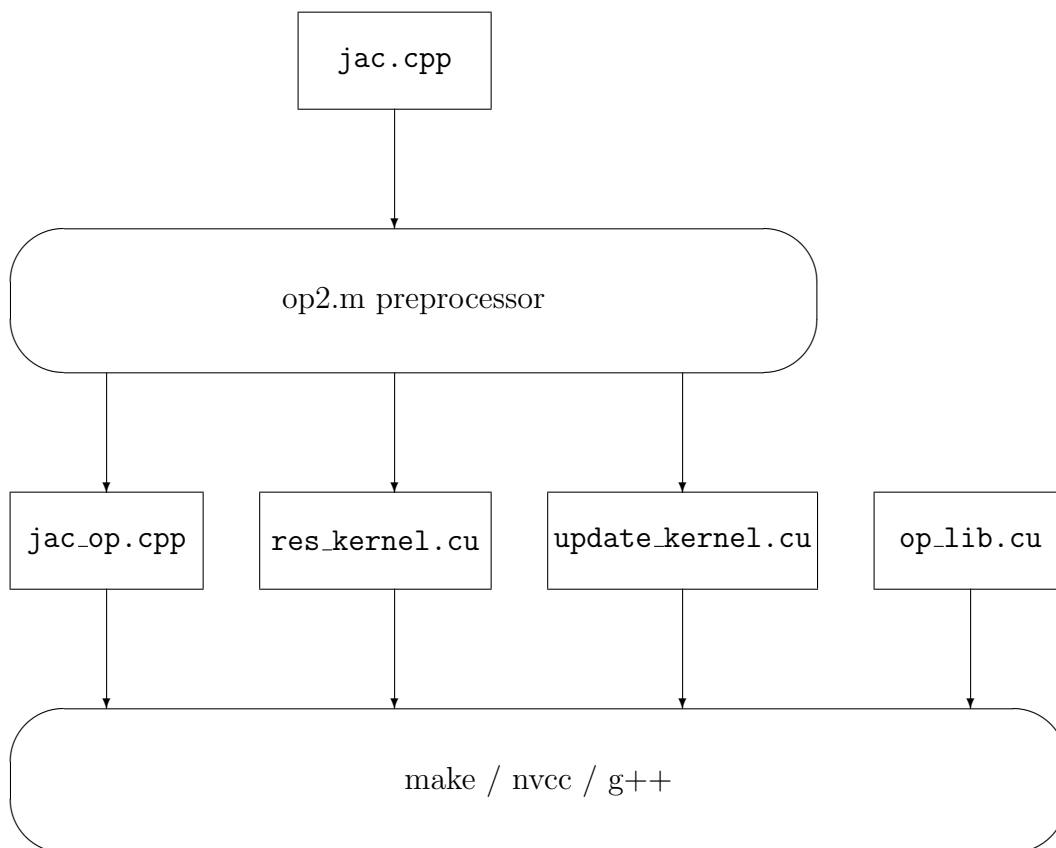Figure 2: Sequential code build process
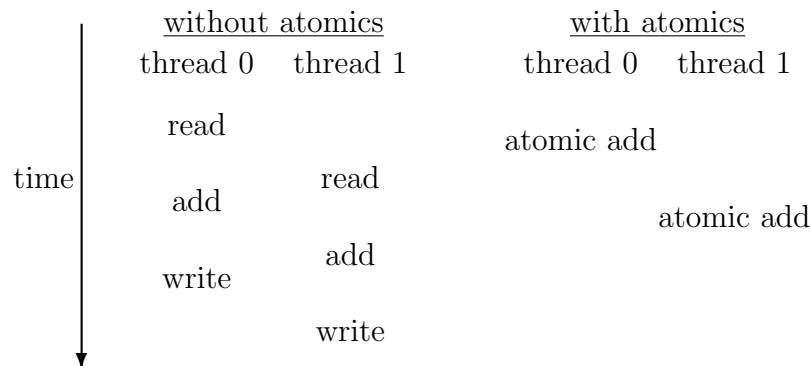

Figure 3: CUDA code build process

	without atomics		with atomics	
	thread 0	thread 1	thread 0	thread 1
read				

Figure 4: Addition with and without atomic execution

12 Atomic operations and global reductions

As illustrated in figure 4, when different threads try to update the same variable at the same time, there is a danger that the resulting answer will be incorrect. In the scenario in the figure, the final value includes the increment from thread 1, but not the increment from thread 0. This problem is avoided by using an atomic add which makes the combined read/add/write a single operation, so that the atomic add by thread 0 is guaranteed to have completed before the start of the atomic add by thread 1.

Atomic operations are supported by the NVIDIA hardware and CUDA, but only for integer operations. However, there is a special atomic compare-and-swap instruction `atomicCAS` which can be used for a software implementation of atomic floating point operations.

The syntax for `atomicCAS` is:

```
int atomicCAS(int* old, int compare, int new);
```

If `old` matches `compare`, the value of `old` is changed to `new`; otherwise it is left unchanged. In either case, the function returns the initial value of `old`. Table 1 shows how this can be used to implement an atomic add. The `float_as_int` and `int_as_float` operations are type casts which do not change the bit-wise values. The key idea is that the routine computes a new value, but this is only stored if the old value hasn't changed in the meantime. If it has, the routine has to repeat the operation until it successfully updates the old value.

This example is easily modified to perform other reduction operations such as computing a minimum or a maximum. This will be useful in performing global reduction operations such as max / min / sum. Each thread in a block can first do the reduction for the set elements it processes. The values for all of the threads in the block can then be combined using the techniques outlined in the CUDA

Table 1: Implementation of floating point atomic addition (code supplied by Jonathan Cohen of NVIDIA Research)

```
float atomicAdd(float *addr, float val) {
  float old = *addr, old2;

  do {
      old2 = old;
      old  = int_as_float(atomicCAS( (int*)addr,
                                     float_as_int(old2),
                                     float_as_int(old2+val) ));
  } while( old2 != old );

  return old;
}
```

SDK `reduction` example, and then the global reduction value can be obtained by thread 0 of eack block performing an atomic update of a single variable in the global memory space. On the Fermi GPU, the global variable will be held in the shared L2 cache and so the `atomicCAS` operation will be performed efficiently with minimal latency.

13   Register pressure

One concern raised by Jamil Appa, based on his experience with porting CFD codes to CUDA, is that each thread can have only a very limited of registers, 128 on the GT200 series and only 64 on Fermi (even though each core has 2048 registers on the GT200 and 1024 on Fermi).

If more than this number are needed then they "spill over" into so-called "local memory" which resides in the main global memory with all of the limitations in latency and bandwidth that entails. With Fermi, this is mitigated only slightly through the use of the 16kB L1 cache, corresponding to 2k DP variables. If there are 128 active threads, this corresponds to just 16 DP variables per thread, so this isn't much of an addition to the register set.

To cope with this, we may need to change the way in which some kernels are implemented. At present, for each element in the active set, the thread copies the required input argument values into registers and then calls the user-supplied element function. For the viscous edge calculation in HYDRA, this would require two 7×4 "gradient" arrays, plus four vectors of size 7 for the flow variables and residuals at the two nodes, and hence at least 84 DP variables, neglecting the registers required to perform the calculations.

For read-only arguments (such as the gradient arrays in the viscous edge calculation) it would be better to simply pass in the pointer to the shared-memory location.

For increment arguments, using shared-memory pointers would raise the problem of data dependencies. When using thread coloring, it would be necessary to wrap it around the entire element calculation, rather than simply the incrementing operation as at present. Because of this, it might be best to re-order the set elements to minimise the number of different colors within each warp. Doing this would disrupt the identity mapping for direct data references, which means an additional mapping would need to be stored and fetched.

Alternatively, the use of atomic updates would either require code transformation of the user-supplied element function, or the user would have to use a special function to perform the increment and this would then be mapped to the required atomic operations.

14 Auto-tuning

The more I get into the details of implementation, the more I see the possibility for different methods of implementation, as well as uncertainty in the optimal values for various parameters which need to be specified. To optimise the run-time performance, I think it will be essential to use auto-tuning techniques, building on ideas already used by packages such as ATLAS and FFTW. To emphasise this point, I don't think I am talking about squeezing out an extra 10% of performance; I think it is more likely that auto-tuning will double the performance.

Here I list some of the choices to be optimised, and the tradeoffs they involve:

- Size of block partition

  To maximise date reuse within eack block, the natural choice is to make the block partition as large as possible subject to the constraint that it fits inside the shared memory. However, making it a bit smaller would allow multiple blocks to run concurrently on the same multiprocessor, and the run-time scheduler would then naturally overlap I/O for one block with computation for another.

- Number of threads per block

  One option is to have one thread per set element, giving the run-time scheduler the maximum freedom. However, this increases the cost of thread synchronisation, might run into difficulties with the total register usage, and doesn't amortise startup costs over multiple elements.

- Thread coloring or atomic increments

  A tradeoff between synchronisation costs and idle threads, versus the cost of hardware-assisted atomic operations.

- Use of registers or shared-memory pointers

  This was explained in the previous section

These choices exist for each parallel loop. What is optimal for one loop may not be optimal for another.

Optimisation strategies which could be used include:

- Brute force

  Specify a small set of possible values for the integer variables (block partition size and number of threads per block) and then systematically try every combination of choices and values.

  This would be particularly applicable if, perhaps through profiling, the user can identify the most important loop(s).

- Genetic algorithms (or other stochastic method)

  Brute force could work at the level of a single loop. To handle lots of loops, one either uses brute force on each one, one at a time, or one could perhaps use a higher-dimensional method like GA to optimise them all at the same time.

- Run-time optimisation

  In some cases, the optimal choice/value may depend on the specifics of the problem being solved, i.e. may depend on the run-time data. If an application requires many timesteps or iterations, it could be feasible at the beginning to try different choices/values for different timesteps/iterations and monitor the execution times to decide on the best combination.

- Machine learning

  Maybe Mike O'Boyle could use his machine learning approach to anticipate the best choice for parallel loop each based on its defining characteristics.

15   OpenMP and SSE/AVX implementation

Intel's Nehalem CPUs have a separate on-chip 256KB L2 cache for each core, and seem to have at least 200GB/s bandwidth from the L2 cache to the core.

Because of this, I think the simplest approach to an OpenMP implementation is to follow a very similar approach to the GPU implementation but with a single thread for each mini-partition. That way, there's no data conflict within a mini-partition, and the OpenMP parallelism would apply to the loop over all of the mini-partitions of a particular colour. The implementation would be simpler than the CUDA implementation since no global→local renumbering would be required, or any of the low-level thread colouring.

A similar strategy could be used for SSE/AVX vectorization as well, with each SSE/AVX vector element handling a different mini-partition. The mini-partitions would need to be reduced in size so that the required number (up to 8 for single precision computation in AVX) can fit into the 256KB L2 cache.

There is a potentially irritating implementation detail concerning the fact that the set size will usually not be an exact multiple of the number of threads times the length of the SSE/AVX vectors. By using recursive bisection, the number of partitions can be made a power of 2, and their size made equal to within a maximum discrepancy of 1. I think it should be possible to colour the mini-partitions so that the number of mini-partitions of each colour is a multiple of the length of the SSE/AVX vectors. The only problem then is that the final SSE/AVX vector for each mini-partition may not be complete, due to the difference in size of the mini-partitions. I think this can be dealt with by replicating one of the final elements to fill the empty bits of the vector. The redundant computations will not affect the final outcome. The same trick can be used to replicate an entire mini-partition if the number of mini-partitions of one colour is not a multiple of the length of the SSE/AVX vectors.

This should result in very clean, efficient OpenMP SSE/AVX code. The complexity will all be hidden in the plan construction.